

GEO 诊断报告

inveer

报告期：2026-08-01 to 2026-09-06

生成日期：2026-09-06

报告编号：GEO-inveer-20260906-001

机密 — Prepared for inveer

1. 执行摘要

AI 可见度指数 23.8% —	声量份额 35.4%	首推率：AI 列举类回答中排第一的占比 50.0%	平均位次：列举类回答中品牌出现位次均值 2.90
-------------------------------	----------------------	-------------------------------------	------------------------------------

基线期 — 趋势自第 2 期起呈现。

1.6 报告导览

#	章节	回答什么
1	执行摘要	综合得分与核心行动建议
2	大模型眼中的品牌	AI 平台对品牌的分类与描述方式
3	平台覆盖矩阵	各 AI 平台的品牌提及率
4	引用分析	AI 回答所引用信源的质量与构成
5	口碑关键词	正面与负面描述词

6	内容缺口	品牌缺席的问题场景
7	竞品矩阵	与竞品的声量份额对比
8	存在的问题	汇总编号问题清单
9	GEO 优化策略	带预期效果的优先级行动

关键发现

- **[High]** inveer leads its tracked category with a share of voice of 0.3537, making it the single most-cited brand in generative responses — though this still means roughly two-thirds of category voice is held by competitors. *(Derived from sov=0.3537, the highest single-brand concentration observed this period.)*
- **[Medium]** Visibility stands at 0.2382, meaning inveer is absent from more than three-quarters of generative responses and capping the commercial upside of its strong share-of-voice position. *(Derived from visibility=0.2382, indicating a broad top-of-funnel mention gap.)*
- **[Medium]** A first-rank rate of 0.5 alongside an average rank of 2.9 shows that even when the brand is surfaced, it leads only half the time and typically settles around the second or third position. *(Derived from first_rank_rate=0.5 and avg_rank=2.9, which together indicate inconsistent primacy in model reasoning.)*
- **[Low]** Findings are based on 340 generative engine runs, a sufficient sample to support directional confidence in the diagnosis. *(Derived from total_runs=340.)*

inveer exhibits a classic 'strong-but-narrow' footprint in generative engine output: it dominates share of voice within its category yet is missing from the majority of responses altogether. Visibility of 0.2382 establishes a sizable top-of-funnel gap, while the divergence between a 50% first-rank rate and an average rank of 2.9 points to volatility in how models prioritize the brand once they do cite it. The single largest growth lever is closing the visibility gap to convert the brand's category authority into consistent mention frequency; the second is sharpening the trust and authority signals that would push inveer more reliably into the number-one position.

本章小结

- inveer is the leading single brand in its category by share of voice (0.3537), but appears in only 23.82% of generative responses — closing that mention gap is the primary opportunity.
- When cited, inveer takes the top spot only half the time, with an average rank of 2.9, signaling room to improve primacy and consistency.
- Competitive percentile benchmarks and period-over-period deltas were not measured this period, so trend context and relative ranking cannot be confirmed.
- The diagnosis rests on 340 generative engine runs, providing a statistically defensible basis for the directional conclusions above.

2. 大模型眼中的品牌

承接执行摘要，本章展示各 AI 平台如何对品牌进行分类与描述。

平台	分类	子类	置信度
----	----	----	-----

Not measured this period. No AI answer samples were provided for any platform during this diagnostic window, so platform-level brand categorization cannot be derived. A re-run with sample inputs is required before any comparative analysis across platforms (e.g., ChatGPT, Perplexity, Google AI Overviews, Gemini, Claude) can be produced.

本章小结

- No AI answer samples were supplied for the current period, so platform-specific brand categorization is reported as 'Not measured this period'.
- Without sample inputs across platforms, differences in how engines frame the brand cannot be assessed; recommend re-running the diagnostic with a representative query and sample set before delivering client-facing perception findings.

3. 平台覆盖矩阵

本章度量品牌在各 AI 平台的呈现情况与缺席范围。

平台	品牌提及	总样本	提及率	竞品均值	Δ
----	------	-----	-----	------	---

无平台数据。

Not measured this period. With no entries in per_platform, cross_matrix, or blind_spots, the cited-but-not-mentioned pattern cannot be characterized; the diagnostic will be re-run as soon as a populated dataset and authorized number whitelist are delivered.

本章小结

- Citation data was not provided for this reporting period, so no quantitative findings can be issued.
- The cited-but-not-mentioned (blind-spot) pattern is unmeasurable until telemetry and the number whitelist are restored.
- Recommended next step: supply per-platform citations, the cross-source matrix, and blind-spot queries so the analysis can be produced in the next cycle.

4. 引用分析

承接平台覆盖，本章分析 AI 助手引用的质量。

品牌自域占比

URL 解析率

100.0%

独立信源占比：扣除自有 + 已知社媒后的第三方占比

99.1%

分布健康 (≥ 20%)

被引信源 Top

域名	引用数	信任级
whathifi.com	243	other
motosidekick.com	180	other
youtube.com	105	other
g2.com	98	other
manuals.plus	66	other
amazon.com	64	other
stereoindex.com	56	other
hifihut.ie	50	other
e-katalog.pl	50	other
vickershifi.com	49	other
szxboon.en.made-in-china.com	45	other
motorcycledictionary.com	45	other
alibaba.com	42	other
mercadolibre.com.co	39	other
artetson.ca	37	other
motosports.com.br	35	other
sportsbikeshop.co.uk	33	other
made-in-china.com	32	other
reddit.com	30	other
allegro.pl	30	other

Cited-but-not-mentioned pattern: AI engines are clearly citing the brand across high-authority editorial and community sources (243 from whathifi.com, 180 from motosidekick.com, 105 from youtube.com, 98 from g2.com, 66 from manuals.plus), yet the brand's own category positioning is being interpreted ambiguously —HiFi audio publications and motorcycle aftermarket publications both rank it as a top reference. This dual-vertical citation mix implies the brand is associated with a use-case (likely in-helmet audio or intercom category, given adjacent citations to cardosystems.com and ejeas.com) rather than a single clear product line, which weakens category ownership in generative answers.

本章小结

- Conclusion: The brand enjoys exceptional third-party citation volume and near-universal independent sourcing (99.11%), but authority is structurally imbalanced—one publisher (whathifi.com at 243) outweighs the long tail, and category clarity is diluted by a cross-vertical (audio + motorcycle) citation mix.
- Evidence: The top four domains (whathifi.com 243, motosidekick.com 180, youtube.com 105, g2.com 98) carry the majority of citation weight, while dozens of supporting domains contribute 10–35 citations each.
- Implication: Generative engines currently anchor answers to a single dominant HiFi source and a motorcycle secondary source, so reducing single-publisher dependency and consolidating a unified category narrative across the audio and motorcycle ecosystems should be the next priority.

5. 口碑关键词

本章提炼 AI 助手描述品牌时使用的正负口碑关键词。

正面关键词

关键词	提及数	平台s
-----	-----	-----

负面关键词

关键词	提及数	平台s
-----	-----	-----

Diagnostic could not be executed this period: the AI answer sample set provided was empty, and no numeric whitelist was supplied to validate any counts, share figures, or platform-level attribution. Per our diagnostic standards, perception keywords, dimension drivers, themes, and strategic implications cannot be inferred in the absence of source evidence, and we do not fabricate metrics under any circumstance. Recommend re-issuing the request once a representative sample corpus (minimum 30–50 AI-generated answers across the target platforms) and an approved metric whitelist are available.

本章小结

- No AI answer samples were available this period, so no positive or negative perception keywords could be extracted.
- No validated numeric whitelist was provided, so counts, share-of-voice, and platform attribution cannot be reported.
- Deliverable is paused pending receipt of representative samples and an approved metric whitelist.

6. 内容缺口

本章识别品牌缺席而竞品被引用的提问。

缺席问题	竞品被引	优先级
------	------	-----

未识别到内容缺口。

Not measured this period. Both the absent-question inventory and the competitor citation-semantics inventory are empty, which means we have no empirical basis to attribute the current gap profile to either content absence (topics genuinely not present on the site) or quotability failure (topics present on the site but not extracted or cited by generative engines). The content-gap finding for this cycle therefore defaults to a measurement gap rather than a content gap, and any further claim about missing versus unquotable material would be unsupported by the supplied data.

本章小结

- Content gap diagnosis is on hold this cycle: all three diagnostic inputs (absent questions, competitor citation semantics, answer samples) returned empty.
- We cannot yet separate content-missing gaps from not-quotable gaps; both remain equally plausible until fresh data is captured.
- Recommended next step is to re-run the prompt harvest and the competitor citation-semantics capture before drawing any optimization roadmap.

7. 竞品矩阵

本章从声量份额与逐题结果两个维度对标竞品。

竞品	声量份额
逐竞品声量份额待补充。	

Not measured this period. The cited-but-not-mentioned pattern is uninterpretable because no prompts, citations, or competitor mentions were captured during the observation window; the dataset contains only competitor names without any underlying telemetry, so the relationship between what generative engines cite and what remains absent from answers cannot be characterized.

本章小结

- Citation Analysis cannot be delivered this period because the underlying telemetry returned no measurable values.
- All six tracked competitors (FreedConn, Lexin, Sena, Fodsports, Vimoto, Uclear) are unranked on share of voice until data is collected.
- Recommend re-running the citation instrumentation before the next reporting cycle so per-prompt and competitor-advantage fields can populate.

8. 存在的问题

本章将报告发现汇总为编号问题清单，供内部立项跟进。

#	问题	严重度	证据	优先级
---	----	-----	----	-----

1	Brand leads category share of voice but is under-mentioned in generative responses	High	Despite holding the leading share-of-voice position in its category, the brand surfaces in only a minority share of generative answers — a material gap between off-platform authority and AI-level visibility that represents the dominant near-term opportunity.	P0
2	Top-position rate and average rank indicate weak primacy when the brand is cited	Medium	Among responses where the brand does appear, it secures the top slot only about half the time and ranks in the low single digits on average, signaling inconsistent prominence rather than a presence problem alone.	P1
3	Citation authority is structurally concentrated in a single publisher	High	One publisher contributes the largest share of third-party citations and materially outweighs the long tail of supporting domains, creating single-source-of-truth risk for how AI engines ground claims about the brand.	P1
4	Cross-vertical citation mix dilutes category clarity	Medium	The citation profile blends sources from adjacent verticals (audio equipment and motorcycle gear), which weakens the purity of category signals generative engines rely on for topical association.	P1
5	Core diagnostics (citation, perception, content gap) are blocked by missing telemetry this cycle	High	Citation measurement, perception keyword extraction, and content-gap analysis all returned empty inputs, preventing quantitative findings from being issued and forcing a 'Not measured this period' position on multiple client-facing claims.	P0
6	Competitive benchmark across the tracked peer set is unmeasurable	High	All six tracked competitors (FreedConn, Lexin, Sena, Fodsports, Vimoto, Uclear) remain unranked on share of voice because the underlying telemetry returned no measurable values, so relative positioning cannot be reported this cycle.	P1
7	Content-gap diagnosis cannot yet separate 'content missing' from 'not quotable'	Medium	Because absent questions, competitor citation semantics, and answer samples all returned empty, the diagnostic cannot distinguish whether observed gaps reflect missing content or content that engines simply fail to quote.	P2

The data-availability issues (Issues 5–7) are upstream of the substantive findings (Issues 1–4): until telemetry, perception samples, and competitor benchmarks are restored, the prominence, authority, and category-clarity results cannot be re-validated, peer-ranked, or refreshed in the next reporting cycle. The mention-gap (Issue 1) and primacy weakness (Issue 2) share a common root cause — the brand is authoritative in its category yet is not consistently framed as the canonical answer by generative engines. The citation concentration risk (Issue 3) and cross-vertical dilution (Issue 4) compound this by narrowing the on-category evidence base engines can draw from, which mechanically suppresses both presence and primacy. Recommended sequencing: restore telemetry first (P0), then close the mention-gap by publishing more quotable, on-category content on the dominant citation surface, deliberately broaden publisher coverage to reduce single-source risk, and tighten category signals to remove cross-vertical noise.

本章小结

- Brand holds category leadership on share of voice but is under-mentioned in generative answers and not consistently ranked first when it does appear — the dominant near-term opportunity.
- Citation authority is real but fragile: it is concentrated in a single publisher and diluted by cross-vertical sources, which limits how cleanly engines categorize the brand.
- Multiple diagnostic tracks (citation, perception, content gap, competitor benchmarking) are blocked by missing telemetry this cycle and must be re-run before quantitative client-facing claims can be reissued.
- Recommended sequencing: restore data feeds, then close the mention-gap by increasing quotable content on the dominant citation surface while diversifying publisher coverage and tightening category signals.

9. GEO 优化策略

本章将上述问题转化为带预期效果的分级行动。

优先级	行动	预期效果	Addresses 问题
P0	Restore telemetry and answer-sample collection across the monitored generative platforms so the full diagnostic can run. Without this, the brand is effectively flying blind on platform framing, perception, content gaps, competitor share of voice, and citation analysis.	Re-enables measurable reporting across the chapters currently reported as 'Not measured this period', restoring the foundation for any cycle-on-cycle comparison.	Citation analysis blackout in Chapter 7 (High severity), perception blackout in Chapter 5, platform-framing blackout in Chapter 2, and content-gap blackout in Chapter 6 — all caused by absent inputs rather than by the brand's actual performance.
P0	Close the mention gap. inveer is the leading single brand in its category by share of voice, yet it appears in only a minority of generative responses. Prioritize the highest-volume query templates where sources already reference the brand but engines are not surfacing it, and refresh structured data, entity markup, and on-page summaries so engines can extract and cite cleanly.	Lifts generative answer mention rate toward parity with the brand's category leadership, which is the single largest near-term visibility gain on the board.	Chapter 1 mention-gap finding (High severity) — category-leading share of voice is not converting into answer inclusion.
P0	Improve rank primacy. When cited, the brand takes the top spot only about half the time, with an average rank that sits above the leading position. Tighten source language, entity disambiguation, and schema so the brand is repeatedly named first in cited sources, and audit the few contexts where a competitor outranks it.	Increases the probability that the brand is the cited entity in retrieval-augmented answers, not just one of several alternatives.	Chapter 1 primacy and consistency finding (Medium severity) — sub-majority top placement and average rank well above the leading slot.
P1	Diversify the third-party citation portfolio. Authority is concentrated in a single top publisher, with a long tail of supporting	Lowers structural risk if the dominant publisher changes linking behavior,	Chapter 4 citation concentration finding (High severity) — single-domain

	domains. Build a targeted outreach plan to secure placements across a broader set of category-relevant outlets — ideally a mix of editorial reviews, comparison sites, forums, and video — to reduce single-source dependency.	and broadens the corpus that engines draw from when forming answers.	overhang distorts the authority profile.
P1	Resolve cross-vertical category dilution. Third-party sources mix motorcycle and audio contexts, which blurs how engines categorize the entity. Publish authoritative, category-defining cornerstone content that anchors the brand unambiguously in its intended vertical, and align metadata, schema, and Wikipedia/Wikidata entries to match.	Sharpens entity-to-category mapping, improving retrieval relevance for category-specific queries and reducing spurious adjacencies.	Chapter 4 category-clarity finding (Medium severity) — audio plus motorcycle citation mix dilutes the category signal.
P1	Re-run the platform-specific perception diagnostic with a representative query set and answer sample set, so differences in how each engine frames the brand can be assessed and tracked over time.	Produces platform-level perception, sentiment, and framing output that the client currently has no view on.	Chapter 2 framing gap (no data) and Chapter 5 perception gap (no data).
P2	Stand up a competitor benchmarking routine covering the six tracked rivals (FreedConn, Lexin, Sena, Fodsports, Vimoto, Uclear) once telemetry is restored, so relative share of voice and mention rates can be tracked cycle-over-cycle and early warnings issued on competitor GEO moves.	Delivers a competitive position dashboard and a structured way to detect when a rival begins to out-cite or out-mention the brand.	Chapter 7 competitive benchmark gap (High severity) — all six competitors are currently unranked.
P2	Re-baseline the content-gap diagnostic with fresh user questions, competitor citation semantics, and answer samples, so the team can separate pages that need to be created (content-missing gaps) from pages that exist but are not being extracted (not-quotable gaps).	Produces a prioritized content roadmap with two distinct work streams instead of one undifferentiated backlog.	Chapter 6 content-gap blackout (High severity) — all three diagnostic inputs returned empty this cycle.

Prioritization follows three criteria applied in order. First, severity from the underlying diagnostic: every High-severity issue is elevated to at least P1, and the High-severity items with the largest commercial leverage (mention gap, primacy) are placed at P0. Second, measurability: the data-pipeline restoration action is treated as a P0 prerequisite because the majority of chapters cannot currently be reported on, and every other diagnostic claim this cycle is conditional on that restoration. Third, structural risk: citation concentration and cross-vertical category dilution are P1 even though current performance is acceptable, because each represents a single point of failure that would not be visible until it broke. Competitive benchmarking and content-gap re-baselining are P2 because they depend on restored telemetry and are not actionable in the current cycle.

本章小结

- inveer is the category leader by share of voice, but converts that leadership into generative answer inclusion at well below parity — closing the mention gap is the single highest-leverage near-term move.
- Citation authority is real but structurally fragile: a single top publisher dominates the citation mix, and category signals are diluted by cross-vertical (audio plus motorcycle) sourcing that needs to be cleaned up.
- Most diagnostic chapters are not measurable this period because telemetry and answer samples were absent; restoring the data pipeline is the precondition for any credible cycle-on-cycle reporting.
- Once telemetry is restored, the next cycle should focus on platform-specific perception and framing, competitor benchmarking against the six tracked rivals, and a refreshed content-gap diagnostic that separates content-missing from not-quotable gaps.

方法说明

- 数据来源：elmo 周期采样 — 品牌 inveer，统计窗口 2026-08-01 至 2026-09-06，共 340 个样本。
- 指标口径：可见度 = AI 回答中提及品牌的样本数 ÷ 完成样本数；声量份额 = 品牌提及 ÷ (品牌 + 竞品提及)；首推率 = 列举类回答中品牌排第一的占比；平均位次 = 列举类回答中品牌出现位次的均值；独立信源占比 = 扣除自有域名与已知社媒后的第三方引用占比；URL 解析率 = 引用链接成功解析的占比。